

Expanding Safe Sets with Learning-Based Barrier Functions and Reachability

Trevor Matthews, Mahesh Kumar A.R., Azra Begzadić, Sander Tonkens, Zhirui Dai, Nikolay Atanasov, Jorge Cortés, and Sylvia Herbert

Abstract—This paper presents GP-CBF, a learning-based framework for synthesizing and continuously adapting a robust control barrier value function for safe robot navigation in unknown environments. The approach first learns a smooth, conservative estimate of the signed distance function from noisy LiDAR data using Gaussian process regression and a Bayesian Hilbert map in real time. This learned signed distance function serves as a candidate control barrier function, which is then formally refined using Hamilton-Jacobi reachability analysis to generate a valid control barrier value function. The resulting value function is updated online as new sensor measurements are added, enabling progressive expansion of the safe set while maintaining probabilistic safety guarantees. The efficacy of our approach is validated through both simulation and hardware experiments on a mobile robot traversing cluttered environments, ensuring safety and outperforming existing baselines.

I. INTRODUCTION

As autonomous systems become integrated into daily life, ensuring their safe operation is critical. Formal methods, such as control barrier functions (CBFs) and Hamilton-Jacobi reachability (HJR), have emerged as powerful tools for enforcing safety. However, these guarantees rely on accurate obstacle boundaries. In many applications, obstacle boundaries are not known a priori and must be constructed from on-board sensor measurements. Since these methods utilize continuous scalar values to represent safe sets, the inaccuracies inherent in sensor reconstruction introduce boundary errors that invalidate the resulting safety certificates.

Signed distance functions. A common environment representation in robotic applications is the signed distance function (SDF), which encodes the distance to the nearest obstacle. Several methods have been proposed to construct SDFs from depth sensors, including volumetric fusion methods [1]–[3] and neural network (NN)-based implicit representations [4]–[6]. However, these methods typically do not quantify uncertainty in the reconstructed obstacle boundary. To address sensing uncertainty, Gaussian Process (GP)-based methods [7]–[9] have been used to learn SDFs while providing uncertainty bounds, yet their computational cost grows rapidly with environment scale.

Control barrier functions. CBFs enforce safety by ensuring the system remains within a safe set defined by a constraint function. However, constructing CBFs is particularly challenging for complex systems under control input bounds and disturbances. As a result, learning-based approaches




Fig. 1: Overview of the proposed learning-based safety framework. Noisy sensor measurements are used to learn a smooth signed distance function using a Bayesian Hilbert Map and Gaussian Process regression. The resulting learned signed distance function serves as a control barrier function candidate, which is refined via reachability analysis into a valid control barrier value function. The resulting value function is then used in an online safety filter to enforce safety constraints on the control input at each time step.

have been proposed to approximate CBFs, most notably using NNs [10]–[12]. GP-based methods have also been used to learn CBFs, yet these methods often impose restrictive assumptions on the system dynamics [13]. Sensor uncertainty has also been addressed through distributionally robust CBF constraints based on Wasserstein ambiguity sets [14].

Hamilton-Jacobi Reachability Analysis. HJR analysis offers a principled framework for ensuring system safety under control bounds and disturbances. This optimal control problem characterizes the set of states from which failure can be avoided, yielding a value function with rigorous safety guarantees for general nonlinear systems. Despite its strengths, this approach encounters two main challenges. First, grid-based dynamic programming solvers suffer from the curse of dimensionality, restricting use to systems with fewer than six dimensions. Second, HJR is commonly implemented as a least-restrictive safety filter, meaning the optimal control is applied only near the safety boundary, which can induce chattering or overly cautious behavior.

Bridging the gap. An initial step to unify the HJR and CBF paradigms is the control barrier value function (CBVF) [15], which leverages HJR-like computations to

construct a valid CBF. This combines the constructiveness of HJR with the efficient online enforcement of CBFs. However, this direct synthesis inherits the high computational cost of HJR analysis, limiting real-time deployment. To address this, the refineCBF framework [16] uses a candidate value function as a “warm-start” for the formal refinement, converging rapidly to a provably safe value function. This framework was extended [17] to adapt to changes in the perceived obstacle boundary in-the-loop. However, this method assumes perfect sensor information and thus may result in safety violations under noisy or incomplete measurements.

Contributions. This paper presents a framework for computing and continuously adapting a measurement-robust CBF from online LiDAR observations. Using the approaches of [8], [13], we first construct a measurement-robust smoothed SDF that serves as a candidate CBF. This candidate is then refined using refineCBF [16] to account for system dynamics and provide probabilistic safety assurances under imperfect sensing. Once initialized, refineCBF runs in-the-loop as the robot navigates unknown environments, maintaining provable safety assurances during navigation.

Specifically, we make the following contributions:

- We use GP regression and Bayesian Hilbert Maps (BHM) to learn a smooth SDF with probabilistic error bounds and show that this safe set is contained within the true safe set with high probability.
- We introduce an in-the-loop refinement of the smooth SDF via HJR, resulting in a valid CBF that accounts for general control-affine robot dynamics. We further provide a technical analysis establishing probabilistic safety guarantees for the resulting filter, while allowing for an expanding safe set.
- We validate this framework in simulation and hardware using a limited-range LiDAR sensor and provide an extensive comparison to five baselines and under varying sensing conditions.

II. BACKGROUND

Consider a control-affine system of the form

$$\dot{x} = f(x) + g(x)u, \quad (1)$$

where $x \in \mathcal{X} \subseteq \mathbb{R}^n$ is the state and $u \in \mathcal{U} \subseteq \mathbb{R}^p$ is the control input. We assume a compact state space $\mathcal{X}$ and a convex and compact input set $\mathcal{U}$. We consider a time horizon $[t, 0]$. For each initial time $t \leq 0$, we denote the set of admissible control signals by $\mathbb{U}(t) := \{u : [t, 0] \rightarrow \mathcal{U} \mid u \text{ is measurable}\}$. We assume the functions $f : \mathcal{X} \rightarrow \mathbb{R}^n$ and $g : \mathcal{X} \rightarrow \mathbb{R}^{n \times p}$ are Lipschitz and bounded. Under these assumptions, given an initial time t , state x , and $u \in \mathbb{U}(t)$, there exists a unique solution $x_{x,t}^u : [t, 0] \rightarrow \mathbb{R}^n$ of (1) with initial condition $x_{x,t}^u(t) = x$. Throughout the paper, the argument t in $x(t)$ and $u(t)$ is omitted for brevity.

A. Hamilton-Jacobi Reachability

HJR is a framework that determines the set of initial states from which a system can reach a target set and/or avoid a failure set. In many robotic applications, we consider a constraint set $\mathcal{L}$, which represents the allowable states of the

system (1). This set is described as a 0-superlevel set of a safety target function $\ell : \mathbb{R}^n \rightarrow \mathbb{R}$, i.e., $\mathcal{L} := \{x \in \mathcal{X} \mid \ell(x) \geq 0\}$. The objective of safe control is to ensure the trajectory stays within constraint set $\mathcal{L}$ for $\tau \in [t, 0]$.

For a given set $\mathcal{L}$, we are interested in computing the viability kernel $\mathcal{S}(t)$ defined as

$$\mathcal{S}(t) := \{x \in \mathcal{X} : \exists u \in \mathbb{U}(t) \text{ s.t. } \forall \tau \in [t, 0], x_{x,t}^u(\tau) \in \mathcal{L}\}, \quad (2)$$

which represents the set of all initial states at time t in the set $\mathcal{L}$ from which an admissible control signal exists that ensures the system will remain in $\mathcal{L}$ over the entire time horizon.

The objective function for this problem is given as $J(x, t, u) = \min_{\tau \in [t, 0]} \ell(x_{x,t}^u(\tau))$, which represents the smallest value of ℓ along the system trajectory over the time horizon. If this value falls below zero, the trajectory has entered the failure set. The associated value function $V(x, t) = \max_{u \in \mathbb{U}} J(x, t, u)$ is the viscosity solution of the following Hamilton–Jacobi–Isaacs variational inequality

$$0 = \min \left\{ \frac{\partial}{\partial t} V(x, t) + H(x, t), \ell(x) - V(x, t) \right\}, \quad (3)$$

$$V(x, 0) = \ell(x),$$

where the Hamiltonian $H : \mathbb{R}^n \times \mathbb{R}^n \rightarrow \mathbb{R}$ is given by $H(x, t) = \max_{u \in \mathcal{U}} \nabla_x V(x, t)^\top (f(x) + g(x)u)$.

Given V , the viability kernel $\mathcal{S}(t)$ for the constraint set $\mathcal{L}$ is $\mathcal{S}(t) := \{x \in \mathbb{R}^n \mid V(x, t) \geq 0\}$. When the viability kernel $\mathcal{S}(t)$ is non-empty, the optimal safety controller $u^*(x, t)$ is given by

$$u^*(x, t) = \arg \max_u \nabla_x V(x, t)^\top (f(x) + g(x)u). \quad (4)$$

B. Control Barrier Functions

We start by formally defining the concept of CBF.

Definition 1 (CBF [18]). *Consider a continuously differentiable function $h : \mathbb{R}^n \rightarrow \mathbb{R}$ and define its superlevel set $\mathcal{C} := \{x \in \mathbb{R}^n \mid h(x) \geq 0\}$. The function h is a CBF of $\mathcal{C}$ for (1) if there exists an extended class $\mathcal{K}$ function α such that, for all $x \in \mathcal{C}$,*

$$\max_{u \in \mathcal{U}} \nabla_x h(x)^\top (f(x) + g(x)u) \geq -\alpha(h(x)). \quad (5)$$

Given a valid CBF h , the safety filter can minimally modify a nominal control law $u_{\text{nom}} : \mathcal{X} \rightarrow \mathbb{R}^p$ to maintain safety using the following quadratic program (QP)

$$\begin{aligned} u^*(x) = \arg \min_{u \in \mathbb{R}^p} & \|u_{\text{nom}}(x) - u\|_2^2 \\ \text{s.t. } & \nabla_x h(x)^\top (f(x) + g(x)u) + \alpha(h(x)) \geq 0. \end{aligned} \quad (6)$$

In practice, it is common to choose $\alpha(x) = \gamma x$, with $\gamma \geq 0$. The synthesis of a CBF for a nonlinear system is challenging, particularly when the system is subjected to bounded control inputs and disturbances. However, recent work has shown HJR can be used to constructively generate CBFs, resulting in control barrier-value functions (CBVFs) defined below.

Definition 2 (CBVF [15]). *A control barrier-value function $h_\gamma : \mathbb{R}^n \times (-\infty, 0] \rightarrow \mathbb{R}$ is given by*

$$h_\gamma(x, t) = \max_{u \in \mathcal{U}_{[t, 0]}} \min_{s \in [t, 0]} e^{\gamma(s-t)} \ell(x_{x,t}^u(s)), \quad (7)$$

for some $\gamma \in \mathbb{R}_{>0}$ and $\forall t \leq 0$, with initial condition $h_\gamma(x, 0) = \ell(x)$.

The CBVF described in (7) can be constructed as in (3). As in [17], we set $\gamma = 0$, while using any $\gamma > 0$ online, such that (7) simplifies to the standard HJR problem. Once the value function h is defined, it can be integrated into a safety filter similar to (6) such that

$$u^*(x) = \arg \min_{u \in \mathbb{R}^p} \|u_{\text{nom}}(x) - u\|_2^2 \quad (8)$$

$$\text{s.t. } \frac{\partial}{\partial t} h(x, t) + \nabla_x h(x, t)^\top (f(x) + g(x)u) + \gamma h(x, t) \geq 0.$$

C. Gaussian Process Regression

GP regression is a non-parametric approach based on the assumption that any finite number of measurements $\{y(x^{(1)}), \dots, y(x^{(N)})\}$, $N \in \mathbb{N}$, of an unknown function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ at inputs $x \in \mathcal{X}$ follows a joint Gaussian distribution [19]. The GP $f(\cdot) \sim \mathcal{GP}(m(\cdot), k(\cdot, \cdot))$ is fully defined by a prior mean $m : \mathcal{X} \rightarrow \mathbb{R}$ and a kernel function $k : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}_{\geq 0}$. By conditioning the prior on training points $(x^{(i)}, y^{(i)})$ with $i = 1, \dots, N$, where $y^{(i)} = f(x^{(i)}) + \epsilon_i$, $\epsilon_i \sim \mathcal{N}(0, \sigma_{\text{on}}^2)$, we obtain a posterior distribution that is again Gaussian, with mean and variance given by

$$\mu(x) = k(x)^\top (K + \sigma_{\text{on}}^2 I_N)^{-1} y, \quad (9a)$$

$$\sigma^2(x) = k(x, x) - k(x)^\top (K + \sigma_{\text{on}}^2 I_N)^{-1} k(x), \quad (9b)$$

with kernel vector $k(x) = [k(x, x^{(1)}), \dots, k(x, x^{(N)})]^\top$, kernel matrix K through $K_{i,j} = k(x^{(i)}, x^{(j)})$, and $y = [y^{(1)} \dots y^{(N)}]^\top$.

III. PROBLEM STATEMENT

Consider a robot with control-affine dynamics (1). We address the problem of controlling the robot to perform tasks while satisfying safety constraints, with a primary focus on obstacle avoidance. For this purpose, we consider an open obstacle set $\mathcal{O} \subset \mathcal{X}$, and its complement constraint set $\mathcal{C} = \mathcal{X} \setminus \mathcal{O}$. We consider the ground-truth SDF $d : \mathcal{X} \rightarrow \mathbb{R}$ of the set $\mathcal{O}$, which measures the signed Euclidean distance from a query point x to the boundary $\partial\mathcal{O}$, with sign indicating whether $x \in \mathcal{O}$, such that

$$d(x) = \begin{cases} \min_{y \in \partial\mathcal{O}} \|x - y\|_2 & \text{if } x \notin \mathcal{O}, \\ -\min_{y \in \partial\mathcal{O}} \|x - y\|_2 & \text{if } x \in \mathcal{O}. \end{cases} \quad (10)$$

We consider a robot equipped with a range sensor that provides discrete noisy boundary measurements of a subset of obstacles $\mathcal{O}$ during deployment. We want to retain safety while obtaining new environment measurements as the robot navigates. Critically, this relies on growing the safe set from its initial estimate, expanding the perceived obstacle-free region. Based on the system description in (1), we address the following.

Problem 1. *Given a robot with dynamics defined in (1) and a range sensor that provides boundary measurements, develop an adaptive safety-filtering framework that ensures safety in environments with unknown static obstacles. The framework should minimally modify nominal control inputs and iteratively reduce conservatism as the environment is mapped, while maintaining formal safety guarantees.*

IV. METHODOLOGY

We develop a method for synthesizing and updating a safety filter online, using noisy distance measurements collected continuously as the system operates over time. Because the true environmental SDF d is unknown, we learn a smooth, robust approximation using a Bayesian Hilbert Map and GP regression. This learned SDF serves as an initial CBF candidate and is refined into a valid CBVF through HJR analysis that accounts for the system dynamics and control bounds, using the SDF as a “warm-start” for the dynamic programming-based computation. The resulting data-driven CBVF is updated online, provides probabilistic safety guarantees, and can be deployed in a real-time safety filter. Continuous learning progressively expands the learned safe set, reducing conservatism as the system improves its estimate of the true safety boundaries.

A. Learning a smooth SDF via BHMs and GP regression

SDFs represent distance to the boundary of obstacle space and hence are a good choice for capturing geometric safety constraints. However, two key elements are required for their downstream use in a safety filter: 1) continuous differentiability (so that a safe control can be computed), and 2) probabilistic outputs (to measure uncertainty about the unknown environment). To achieve these properties, we learn a smooth surrogate function $\tilde{d}$ of the true SDF d using GP regression. In our online setting, the robot gathers and updates surface boundary measurements during operation rather than relying on a pre-collected dataset. To do this, we employ BHMs [20] to obtain a dataset of surface points by aggregating boundary measurements over time and inferring occupancy across the environment. Then, we run marching cubes [21] on the BHM to extract surface points p_i .

To train a GP SDF model on this data, we follow the Log-Gaussian Process Implicit Surface approach in [8], which leverages Varadhan’s distance formula [22] to obtain an unsigned distance function. We consider a GP $s(\cdot) \sim \mathcal{GP}(0, k_s(\cdot, \cdot))$ with zero prior mean and a Matérn 3/2 kernel with kernel scale $\sqrt{3}/\lambda$. This kernel scale and shape are chosen to model the heat diffusion process underlying Varadhan’s formula. Then, the unsigned distance $\tilde{s}(x)$ at a query point x is obtained by taking the logarithm of the GP inference, resulting in $\tilde{s}(x) = r(s(x)) = -\log(s(x))/\lambda$, which inverts the kernel’s exponential “diffusion” to recover a linear metric distance. We denote $r(z) = -\log(z)/\lambda$ as the nonlinear transformation function. Based on this transformation, we construct our training dataset $\mathcal{D} = \{(p_i, 1)\}_{i=1}^{N_k}$ for $s(\cdot)$ by associating each surface point p_i with a target measurement of $y_i = 1$. This ensures that $r(1) = 0$, correctly mapping surface points to zero distance.

To obtain guarantees on the behavior of this function, we make the following assumption on the function s .

Assumption 1. *The function s lies in a Reproducing Kernel Hilbert Space (RKHS) associated with a kernel $k_s(\cdot, \cdot)$ such that s has a bounded RKHS norm, i.e., $\|s(x)\|_{\mathcal{H}_{k_s}} \leq B$.*

This assumption allows us to establish the following prediction error bound for GP regression.

Lemma 1. ([23]) *Under Assumption 1, for a given $\delta \in (0, 1)$, there exists a constant $\beta > 0$ such that the following holds on a compact set $\mathcal{X}$*

$$\Pr\{|\mu_s(x) - s(x)| \leq \beta\sigma_s(x), \forall x \in \mathcal{X}\} \geq 1 - \delta. \quad (11)$$

Determining the constant β in Lemma 1 typically relies on prior knowledge of the underlying function and the choice of kernel parameters [23]. While this aspect is important for practical implementation, a detailed treatment of selecting β lies beyond the scope of this paper. Since the nonlinear transformation $r(z) = -\log(z)/\lambda$ is monotonically decreasing, applying it to both sides of the inequality in (11) flips the signs. Hence, we have that

$$r(\mu_s(x) + \beta\sigma_s(x)) \leq \tilde{s}(x) \leq r(\mu_s(x) - \beta\sigma_s(x)) \quad (12)$$

holds with a probability of at least $1 - \delta$. With the bounds on the unsigned distance established, we now address the sign estimation to obtain the SDF.

Note that so far, our distance estimate $\tilde{s}(x)$ is unsigned. To obtain the sign, we employ the previously mentioned BHM. Beyond using it to obtain the dataset of surface points $\mathcal{D}$, the BHM provides a probabilistic estimate $\tilde{b}(x)$ of the occupancy of each state $b(x)$. We denote $b(x) = 1$ if x is free, i.e., $x \in \mathcal{O}^C$, and $b(x) = -1$ if $x \in \mathcal{O}$. The estimate $\tilde{b}(x)$ has probability $\Pr(b(x) = 1) = q$ and $\Pr(b(x) = -1) = 1 - q$. Hence, our SDF estimation becomes:

$$\tilde{d}(x) = \tilde{b}(x)\tilde{s}(x). \quad (13)$$

Then, $\tilde{d}(x)$ has a mixture distribution with density:

$$p(\tilde{d}(x)) = \begin{cases} qp_{\tilde{s}}(\tilde{s}(x)) & \tilde{d}(x) \geq 0, \\ (1-q)p_{\tilde{s}}(\tilde{s}(x)) & \tilde{d}(x) < 0, \end{cases} \quad (14)$$

where $p_{\tilde{s}}(\tilde{s}(x))$ is the density of $\tilde{s}(x)$. This gives $\mathbb{E}[\tilde{d}(x)] = q\mathbb{E}[\tilde{s}(x)] + (1-q)\mathbb{E}[-\tilde{s}(x)] = (2q-1)\mathbb{E}[\tilde{s}(x)]$ and $\mathbb{V}[\tilde{d}(x)] = 4q(1-q)\mathbb{E}^2[\tilde{s}(x)] + \mathbb{V}[\tilde{s}(x)]$.

When x is in the free space, $q \rightarrow 1$; when x is in occupied space, $q \rightarrow 0$. Both cases give $\mathbb{E}[\tilde{d}] = \tilde{b}(x)\mathbb{E}[\tilde{s}(x)]$ and $\mathbb{V}[\tilde{d}(x)] = \mathbb{V}[\tilde{s}(x)]$. Assuming a well-trained BHM, q quickly transitions between 0 and 1 across the surface, hence, $\mathbb{V}[\tilde{d}(x)] \approx \mathbb{V}[\tilde{s}(x)]$, which means we can simply use a deterministic sign prediction with BHM, i.e., $b(x) = 1$ if $q \geq \tau$, or $b(x) = -1$ if $q < \tau$, where $\tau \in (0, 1)$ is the decision boundary. Although $\tau = 0.5$ by default, we can select a larger τ to ensure the resulting $\tilde{d}(x)$ is conservative, by requiring high probability τ for a state x to have $\tilde{d}(x) > 0$.

Since $\tilde{d}(x)$ has a mixture distribution, we utilize a deterministic sign prediction from the BHM to obtain a probabilistic pessimistic estimate of the SDF. Therefore, given the decision boundary $\tau \in (0, 1)$, for some small constant $\epsilon > 0$ we define our conservative estimate $\underline{d}(x)$ as

$$\underline{d}(x) = \begin{cases} r(\mu_s(x) + \beta\sigma_s(x)), & \text{if } q \geq \tau, \\ -r(\max(\epsilon, \mu_s(x) - \beta\sigma_s(x))), & \text{if } q < \tau, \end{cases} \quad (15)$$

where ϵ ensures the argument of the logarithm remains strictly positive in regions of high GP uncertainty. This piecewise formulation guarantees that $\underline{d}(x) \leq \tilde{d}(x)$ for all x with a probability of at least $1 - \delta$.

However, $r(z)$ distorts the variance. Therefore, in practice, we approximate $\mathbb{V}[\tilde{s}(x)]$ by a softmax fusion process as in [8]. With the aforementioned dataset of surface points $\mathcal{D}$ and associated variances, we can write

$$\tilde{s}(x) \approx m(x, \{p_i\}_{i=1}^N) = \sum_i w_i z_i, \quad (16)$$

with $z_i = \|p_i - x\|_2$, $w_i = \frac{\exp(-\alpha z_i)}{\sum_{i=1}^N \exp(-\alpha z_i)}$, $\alpha > 0$, and m the softmax value. Then, we can calculate the approximation of $\mathbb{V}[\tilde{s}(x)]$ with the first-order approximation of (16):

$$\mathbb{V}[\tilde{s}(x)] \approx \sum_{i=1}^N \|\nabla_{p_i} m(x)\|_2^2 \sigma_i^2, \quad (17)$$

with $\mathbb{V}[\tilde{d}(x)] = \mathbb{V}[\tilde{s}(x)]$ as before.

B. CBVF construction

Since directly using the posterior mean of the SDF surrogate $\mu_{\tilde{d}}$ may overestimate safety in uncertain regions, potentially leading to labeling unsafe states as safe, we consider the pessimistic estimate of the SDF $\underline{d}$ from Eq. (15). If the kernel parameters are suitably chosen, the resulting superlevel set of $\underline{d}$ is an inner approximation of the real safe set $\mathcal{C}$ as formalized in the following result.

Lemma 2. *Consider a prior $\mathcal{GP}(0, k_s(\cdot, \cdot))$ with kernel k_s . Given Assumption 1 and training dataset $\mathcal{D}$, the 0-superlevel set of the GP-SDF $\underline{d}$, $\mathcal{C}_s = \{x \in \mathcal{X} : \underline{d}(x) \geq 0\}$, is included in $\mathcal{C}$ with a probability of at least $1 - \delta$, i.e., $\mathbb{P}(\mathcal{C}_s \subseteq \mathcal{C}) \geq 1 - \delta$.*

The proof of this result follows an analogous argument to [13, Lemma 5.5] and we omit it for space reasons. The resulting GP-SDF $\underline{d}$ can now be used as the constraint function in an HJR formulation to construct and refine a CBF that is valid with respect to the general control-affine dynamics (1) with bounded control inputs. refineCBF uses a discretized version of (3) for its update:

$$h^{(k+1)}(x) \leftarrow \min \left\{ \ell_j(x), h^{(k)}(x) + \Delta^{(k)} L_f^* h^{(k)}(x) \right\}, \quad (18)$$

where $\ell_j(x) = \underline{d}_j$ and $\Delta^{(k)}$ is the time step.

At initialization, both the constraint function and the initial approximate CBF are set as the initial pessimistic SDF, i.e., $\ell_0(x) = h^{(0)}(x) = \underline{d}_0(x)$. At the initial iteration, our method runs (3) until convergence, i.e., $|h_0^{(k)}(x) - h_0^{(k-1)}(x)| < \epsilon$ for $\epsilon \in \mathbb{R}_{>0}$, resulting in the converged CBF $h_0^*(x)$. After initial convergence of h_0^* , we update the obstacle function $\ell(x)$ through updates to the GP node and the CBF by applying (18) in parallel, as seen in Algorithm 1. The CBF update is based on the latest available $\ell(x)$.

The theoretical guarantee for this refinement is provided by the following result, which ensures convergence to a provably safe CBVF.

Theorem 1. (Convergence to a safe CBVF [17], Thm 1) *If the value function (18) converges to $h^*(x)$ as $t \rightarrow -\infty$, then $h^*(x)$ is a safe CBVF for all $x \in \mathcal{H}^*$, its 0-superlevel set.*

However, the guarantees provided by Theorem 1 are difficult to apply directly in an online setting for two main

Algorithm 1 Asynchronous GP-SDF and CBVF Update

```

1: Initialize  $d_0$ , set  $\ell_0(x) \leftarrow \underline{d}_0(x)$ , and  $h_0(x) \leftarrow \underline{d}_0(x)$ 
2:  $h_0^*(x) \leftarrow$  Solve (18) until convergence
3:  $k \leftarrow 1$ ,  $j \leftarrow 1$ ,  $h^{(0)}(x) \leftarrow h_0^*(x)$ 
   GP-SDF (runs continuously at rate  $f_{\text{GP}} = \frac{1}{\Delta_{\text{GP}}}$ )
4: loop
5:   Update measurements  $\mathcal{D}_j = \mathcal{D}_{t_k - \Delta_{\text{GP}}} \cup \dots \cup \mathcal{D}_{t_k}$ 
6:   Update GP-SDF  $d_j(x)$  given  $d_{j-1}$  and  $\mathcal{D}_j$ 
7:    $j \leftarrow j + 1$ 
   refineCBF (runs continuously at rate  $f_{\text{CBF}} = \frac{1}{\Delta_{\text{CBF}}}$ )
8: loop
9:   Query current constraint function  $\ell_k(x) = \underline{d}_j(x)$ 
10:   $h^{(k+1)}(x) \leftarrow \min\{\ell_k(x), h^{(k)}(x) + \Delta L_f^* h^{(k)}(x)\}$ 
11:   $k \leftarrow k + 1$ 
   Controller (runs continuously at rate  $f_{\text{ctrl}}$ )
12: loop
13:   Query CBF  $h = h^{(k)}(x)$ 
14:   Apply  $u \leftarrow u^*(x)$  via the safety-filter QP (20)

```

reasons. First, Theorem 1 assumes a deterministic representation of the safe set $\mathcal{C}$. Second, its guarantees are only valid upon convergence, which would require running HJR until convergence upon each SDF estimate update.

In contrast to Theorem 1, we instead have an initial candidate CBF which is based on GP error bounds, providing only a probabilistic guarantee of set inclusion. Next, we provide the following result, which extends Theorem 1 and the deterministic refinement framework of [17] to include formal probabilistic guarantees under a learned, uncertain environment model. With this result, we can use a GP-SDF-based candidate CBF as an initial guess for a dynamic programming refinement framework, and obtain a probabilistically safe value function for our system as a result.

Theorem 2. *Let h_s be a learned candidate CBF function and let its 0-superlevel set be $\mathcal{C}_s = \{x \in \mathcal{X} : h_s(x) \geq 0\}$ for system (1). Assume that $\mathcal{C}_s$ provides a conservative approximation of the true, unknown constraint set $\mathcal{C}$, meaning it is a subset of the true set with high probability, i.e., $\mathbb{P}(\mathcal{C}_s \subseteq \mathcal{C}) \geq 1 - \delta$. If the value function initialized with the initial guess $h_0(x) = h_s(x)$ in (18) converges to a function h^* with 0-superlevel set $\mathcal{H}^*$, then*

- 1) $\mathcal{H}^*$ is a control invariant set,
- 2) $\mathcal{H}^*$ is a subset of the true set $\mathcal{C}$ with high probability, i.e., $\mathbb{P}(\mathcal{H}^* \subseteq \mathcal{C}) \geq 1 - \delta$.

Proof. Let h^* be the CBVF and $\mathcal{H}^* := \{x \in \mathcal{X} \mid h^*(x) \geq 0\}$ due to Theorem 1. By construction, h^* is the viscosity solution of the HJI (3) given the system description in (1) with convex and compact $\mathcal{U}$. For each $x \in \mathcal{H}^*$, the HJI implies the existence of an optimal safety controller u^* such that $\nabla h^* \cdot f(x, u^*) \geq 0$. This control law guarantees that on the boundary, i.e., $h^*(x) = 0$, the vector field points inside. Hence, trajectories starting in $\mathcal{H}^*$ do not leave the set under u^* . By Nagumo's theorem, the set $\mathcal{H}^* = \{x \mid h^*(x) \geq 0\}$ is forward invariant under the closed-loop dynamics (1), which proves the first claim. In order to prove the second claim, we recall the assumption on the inclusion of the sets,

i.e., $\mathbb{P}(\mathcal{C}_s \subseteq \mathcal{C}) \geq 1 - \delta$. For $h(x, t)$ defined through (18) and any finite horizon $t < 0$, the refined set is defined as $\mathcal{H}_t = \{x \in \mathcal{X} : h(x, t) \geq 0\}$, which consists of all states from which there exists a control policy whose trajectory remains in $\mathcal{C}_s$ and terminates in $\mathcal{C}$. Since $\mathcal{C}_s \subseteq \mathcal{C}$ with probability of at least $1 - \delta$, the terminal state is in $\mathcal{C}$ with the same probability. Hence, $\mathcal{H}_t \subseteq \mathcal{C}$ with probability of at least $1 - \delta$ for all $t < 0$. Hence, we have

$$\mathbb{P}(\mathcal{H}_t \subseteq \mathcal{C}) \geq \mathbb{P}(\mathcal{C}_s \subseteq \mathcal{C}) \geq 1 - \delta \quad (19)$$

for all $t < 0$. If h converges to h^* as $t \rightarrow -\infty$, the property extends to the infinite-horizon safe set $\mathcal{H}^*$, i.e., $\mathbb{P}(\mathcal{H}^* \subseteq \mathcal{C}) \geq 1 - \delta$, concluding the proof. $\square$

Secondly, to address the issue of requiring convergence of (18), we introduce the following:

Lemma 3. *(Refining a safe CBVF will preserve safety [17, Lemma 4]) If there exists $t_1 \leq 0$ such that $h(x, t_1)$ from (18) is control invariant, $h(x, t)$ is control invariant for all $t < t_1 \leq 0$.*

Critically, this provides a guarantee of preserving safety while converging. By inspection of (18), Lemma 3 holds when we obtain further observations of the environment under the following assumption:

Assumption 2. *(Growing conservative safe set)¹ We assume that with each update of the GP $d(x)$, the safe set grows; in particular, we have $\ell_{j+1}(x) \geq \ell_j(x)$ for all $x \in \{x \mid \ell_{j+1}(x) \geq 0\}$, with $\ell_j(x) = \underline{d}_j(x)$, while remaining a subset of the true obstacle-free region with high probability, i.e., $\mathbb{P}(\{x \mid \ell_j(x) \geq 0\} \subseteq \mathcal{C}) \geq 1 - \delta$ at every iteration j .*

Lemma 4. *(Refining a safe probabilistic CBVF preserves probabilistic safety) Consider Alg. 1. Given Assumption 2, Alg. 1's value function is a probabilistically valid CBVF at every iteration as its 0-superlevel set $\mathcal{C}_s^{(k)}$ is a subset of the true set with high probability, i.e., $\mathbb{P}(\mathcal{C}_s^{(k)} \subseteq \mathcal{C}) \geq 1 - \delta$.*

Proof. Theorem 2 guarantees that $h^{(0)}(x)$ is control invariant and its 0-superlevel set is a subset of $\mathcal{C}$. Then, Lemma 3 ensures $h^{(k)}$ is control invariant and satisfies the CBF constraint for all $x \in \mathcal{C}_s^{(k)} = \{x \mid h^{(k)}(x) \geq 0\}$. Lastly, Eq. (18) guarantees $h^{(k+1)}(x) \leq \ell_j(x)$, thus $\mathcal{C}_s^{(k+1)} \subseteq \{x \mid \ell_j(x) \geq 0\}$, which by Assumption 2 is a subset of $\mathcal{C}$ with high probability. $\square$

This is crucial for navigation in a priori unknown static environments, where we continuously improve our guess of the true SDF in the environment. This assumption can typically be satisfied by characterizing the GP-SDF such that unobserved areas of the state space have a negative value, i.e., are considered unsafe. To improve notational clarity, we omit the iteration argument j .

This result extends the deterministic refinement framework to handle uncertainty, guaranteeing if an initial safe set lies

¹Assumption 2 is reasonable for static environments, where new measurements primarily reveal previously unobserved free space and expand the conservative safe-set estimate. The current framework does not account for substantial safe-set contractions arising from environmental changes; extending the analysis to dynamic environments is left for future work.

within the true non-obstacle set with high probability, then the safe set characterized by the converged value function is included in the non-obstacle set with the same probability. Moreover, the safe set can be expanded upon new observations, enabling exploration of unknown environments.

C. Safe control synthesis

Assuming we remain stationary at the initial point until (18) converges, as in Alg. 1, any future CBVF $h^{(k)}$ for $k > 0$ characterizes a control invariant set (see Lemma 3). Therefore, the value function $h^{(k)}$ serves as a valid CBVF for the system (1). This function can be integrated into a real-time safety filter, which minimally modifies a desired nominal control law as given in (6) such that

$$\begin{aligned} u^*(x) &= \arg \min_{u \in \mathbb{R}^p} \|u_{\text{nom}}(x) - u\|_2^2 \\ \text{s.t. } \nabla_x h^{(k)}(x)^\top (f(x) + g(x)u) + \alpha(h^{(k)}(x)) &\geq 0. \end{aligned} \quad (20)$$

The existence of a solution to this QP is crucial for guaranteeing safety as formalized in the following result.

Proposition 1. *Given the value function $h^{(k)}$ for the control-affine system (1), the optimization problem (20) is feasible for $x \in \{x \mid h^{(k)}(x) \geq 0\}$ whenever the gradient $\nabla_x h^{(k)}(x)$ exists. Furthermore, any solution u^* from (20) renders its 0-superlevel set forward invariant.*

Proof. The proof follows the argument of [15, Proposition 3]. Since $h^{(k)}$ is the value function of (18), it satisfies the CBVF variational inequality almost everywhere [15, Remark 1]. If $h^{(k)}(x) < \ell(x)$, the variational inequality implies $\max_{u \in \mathcal{U}} \nabla h^{(k)}(x)^\top (f(x) + g(x)u) + \alpha(h^{(k)}(x)) = 0$, whereas if $h^{(k)}(x) = \ell(x)$, the same quantity is non-negative. Hence, there always exists $u \in \mathcal{U}$ satisfying the CBF constraint, proving feasibility of (20). For any optimal solution $u^*(x)$, we have $\dot{h}^{(k)}(x) \geq -\alpha(h^{(k)}(x))$. By the Comparison Lemma, $h^{(k)}(x(t)) \geq 0$ whenever $h^{(k)}(x(0)) \geq 0$, implying $\{x \mid h^{(k)}(x) \geq 0\}$ is forward invariant. $\square$

V. EVALUATION

Robot Dynamics. We evaluate the method in simulation and hardware² on a Clearpath Jackal modeled as a nonholonomic system with state $x = [p_x, p_y, \theta, v]^\top \in \mathbb{R}^2 \times [-\pi, \pi) \times [0, 1]$ and control input $u = [a, \omega]^\top \in \mathcal{U}$, governed by

$$\dot{p}_x = v \cos(\theta), \quad \dot{p}_y = v \sin(\theta), \quad \dot{\theta} = \omega, \quad \dot{v} = a. \quad (21)$$

Here, p_x and p_y denote position in a 2-D Cartesian coordinate system, θ the orientation, and v the linear velocity. To reflect actuator limits, we constrain the inputs to $a \in [-0.5, 0.5] \text{ m/s}^2$ and $\omega \in [-1.5, 1.5] \text{ rad/s}$.

Objective and Environments. In each experiment, a nominal, safety-agnostic PD controller drives the robot toward a sequence of goals in an a priori unknown environment, while the respective safety filter modifies the control input. We consider two environments: SimpleNav (a constrained environment with 4–5 obstacles), and Warehouse (a factory environment with objects of varying shapes and sizes). The

safety constraints enforce obstacle avoidance and a minimum forward velocity $v_{\min} = 0.1 \text{ m/s}$ for deadlock mitigation. Performance is evaluated by the number of goals reached, elapsed time until completion or failure, distance traveled, and the minimum safety value along the executed trajectory, where a value of 0.00 denotes a collision and safety violation.

GP Parameters. The surrogate function $s(x)$ is modeled using a GP with a Matérn kernel of order $\nu = 3/2$, trained online using a maximum of 256 measurements at each iteration. The kernel’s length scale is $l = 0.087$, with a characteristic length scale $\lambda = 20$. To ensure real-time performance, we use a sparse approximation with 100 basis functions arranged on a 10×10 grid. The GP operates over a bounded local spatial domain of half-width 2.5 m in each coordinate, enabling efficient localized updates of the SDF.

Baselines. To evaluate the performance of our approach, we consider the following comparisons:

- 1) Oracle: CBVF computed online via HJR leveraging an offline mapped SDF (known environment, upper bound on achievable performance).
- 2) GP-SDF: Safety filter using the GP-learned SDF directly, without value-function refinement.
- 3) SDF-CBF (No GP): Our method using an SDF computed from a distance transform of a discretized occupancy map rather than a GP.
- 4) D* Lite [24]: Path planning method updated with the GP-SDF; if the goal is unobserved, a flood-fill algorithm selects the nearest reachable free-space target.
- 5) iSDF [6]: Our method using a continually learned neural SDF rather than a GP-SDF.
- 6) GP-CBF (Ours): Our proposed method, combining the GP-SDF with online CBVF refinement via refineCBF.

A. Simulation experiments

In simulation, SimpleNav consists of a $10 \text{ m} \times 10 \text{ m}$ room with four cylindrical obstacles, and the Warehouse is a $30 \text{ m} \times 50 \text{ m}$ warehouse with narrow corridors formed by shelving and support pillars. The simulated LiDAR has a 30 m range and additive zero-mean Gaussian noise $\mathcal{N}(0, \sigma^2)$ with standard deviation $\sigma = 0.05$.

Table I summarizes quantitative performance, and Fig. 2 shows representative trajectories and learned safe sets. The Oracle and GP-CBF are the only methods that reach all goals in both environments. SDF-CBF (No GP) constructs the SDF directly from noisy measurements without confidence bounds, overestimating clearance near obstacles and resulting in a collision. In contrast, our method synthesizes a conservative boundary estimate, ensuring the learned safe set remains contained within the true safe set. GP-SDF learns a smooth SDF, but applies it directly as a CBF without incorporating system dynamics or control limits, preventing guarantees of control invariance. D* Lite fails when the goal lies beyond the observed obstacle boundary, as the planner is unable to infer a feasible route through unobserved areas. In contrast, iSDF attempts to infer a global map from local observations, but sensor noise reduces the SDF’s accuracy too early in the learning process to maintain safety.

Our method requires about 32% more time than the Oracle in SimpleNav, and about 100% more time in the

²The implementations of the GP-SDF mapping algorithm and HJR refinement are available on GitHub: [GP-SDF Mapping](#) and [RefineCBF](#).

TABLE I: Performance comparison of control policies in the simulated SimpleNav and Warehouse environments.

SimpleNav Environment				
Control Policy	Goals Reached	Elapsed Time (s)	Distance Traveled (m)	Min Safety Value
Oracle	5/5	125	44	0.35
GP-SDF	4/5	56	30	0.00
SDF-CBF	5/5	142	64	0.01
D* Lite	5/5	96	28	0.36
iSDF	5/5	134	44	0.26
GP-CBF (Ours)	5/5	165	63	0.12

Warehouse Environment				
Control Policy	Goals Reached	Elapsed Time (s)	Distance Traveled (m)	Min Safety Value
Oracle	9/9	896	262	0.25
GP-SDF	0/9	7	2	0.00
SDF-CBF	2/9	178	55	0.00
D* Lite	2/9	98	57	0.00
iSDF	0/9	43	12.5	0.00
GP-CBF (Ours)	9/9	1780	422	0.05

TABLE II: Performance of GP-CBF under varying sensing range and measurement noise in the Warehouse simulation.

Configuration (range, noise)	Goals Reached	Elapsed Time (s)	Distance Traveled (m)
30 m, $\sigma = 0.01$	7/7	1120	284
30 m, $\sigma = 0.05$	7/7	1780	422
15 m, $\sigma = 0.03$	7/7	1701	580
8 m, $\sigma = 0.01$	0/7	48	10
8 m, $\sigma = 0.05$	0/7	47	7

Warehouse. This arises from narrow passages and partial observability: conservative initialization treats unobserved areas as unsafe, so the robot must detour and revisit regions before certifying safe passage. Despite this increased travel time, Fig. 3 shows the GP-SDF and associated CBVF progressively expanding toward the Oracle safe set as LiDAR measurements reveal previously unobserved safe regions.

In our implementation, HJR refinement runs at approximately 3 Hz, while the GP-SDF updates at 10 Hz in SimpleNav and 2 Hz in Warehouse. The slower Warehouse rate is primarily due to the cost of querying the SDF over the discretized HJR grid. Because GP updates operate over a fixed local domain, their cost depends mainly on the number of basis functions rather than the overall environment size. A 40 Hz controller asynchronously enforces safety using the latest CBVF, so slower updates primarily reduce exploration efficiency rather than safety, producing more conservative behavior such as revisiting explored regions before newly discovered paths become available.

To assess robustness, we vary LiDAR range and measurement-noise standard deviation over seven Warehouse goals (Table II). Increased noise raises elapsed time by making the safe set more conservative, but does not reduce the completion rate. In contrast, severely limited sensing leaves much of the environment uncertified as safe, confining the robot near its initial position and preventing progress.




Fig. 2: Trajectories under different control policies in the Warehouse simulation. Each subplot shows the robot trajectory after reaching goals 1, 2, 3, and 9 (top left to bottom right), along with the corresponding safe regions estimated by GP-CBF.




Fig. 3: Evolution of the safety value and the fraction of safe states in the Warehouse simulation. **Top:** Safety value along the executed trajectory. **Bottom:** Fraction of the state space certified as safe.

B. Hardware experiments

Complementing the simulation results, we carry out real-world experiments using a Clearpath Jackal. On board, the robot is equipped with an Intel Core i7-9700TE CPU with 16 GB RAM, an NVIDIA GeForce GTX 1650 GPU, and an Ouster OS1-32 LiDAR sensor. The SimpleNav environment is a confined area with five rectangular obstacles, and the Warehouse is a mock factory floor, shown in Fig. 1.

Results are summarized in Table III. As in simulation, GP-SDF performs poorly, failing to reach goals that require safety filtering beyond the nominal controller. In contrast, GP-CBF achieves comparable performance to the Oracle in both elapsed time and distance traveled. The stronger relative performance compared to the Warehouse simulation is due to the smaller scale of the real environments, allowing the GP-CBF to capture relevant geometry with fewer exploration steps, as reflected in Fig. 4. Although the fraction of safe states initially increases, it later levels off and fluctuates as the refinement evolves. This is due to increased model uncertainty and measurement noise in hardware, yielding more conservative predictions in sparsely observed regions.

VI. CONCLUSIONS

We present a learning-based framework that constructs a continuously differentiable GP-based SDF and refines

TABLE III: Performance comparison of control policies in hardware for the SimpleNav and Warehouse environments.

SimpleNav Environment				
Control Policy	Goals Reached	Elapsed Time (s)	Distance Traveled (m)	Min Safety Value
Oracle	6/6	114	23	0.02
GP-SDF	3/6	58	13	0.00
GP-CBF (Ours)	6/6	93	23	0.05

Warehouse Environment				
Control Policy	Goals Reached	Elapsed Time (s)	Distance Traveled (m)	Min Safety Value
Oracle	5/5	94	31	0.13
GP-SDF	2/5	47	15	0.00
GP-CBF (Ours)	5/5	104	34	0.05

Figure 4 consists of three vertically stacked plots. The top plot shows the Warehouse environment map with five goals (0-5) and obstacle boundaries. It displays trajectories for GP-CBF (Ours) in blue, GP-SDF in yellow, and Oracle in pink. The middle plot shows the Safety Value (y-axis, -1 to 2) over time (x-axis, 0 to 100s) for the same three methods. The bottom plot shows the fraction of Safe Cells (y-axis, 0.4 to 0.6) over time (x-axis, 0 to 100s). A legend at the bottom identifies the lines for GP-CBF (Ours), GP-SDF, Oracle, Goals (0-5), Crash, Obstacle boundary, and CBF Safe region.

Fig. 4: Performance comparison in the Warehouse environment. **Top:** Robot trajectories for reaching five selected goals under different control methods. **Middle:** Evolution of the value function along the robot's trajectory for each method. **Bottom:** Fraction of the state space certified as safe over time.

it into a valid CBVF. The method provides probabilistic safety guarantees under sensor uncertainty and is validated in simulation and hardware across diverse environments and noise conditions. The current framework is limited to static environments and low-dimensional systems due to the HJR refinement step. Future work will address dynamic environments, alternative sensing modalities, and scalable extensions to higher-dimensional systems such as quadcopters.

REFERENCES

- [1] H. Oleynikova, Z. Taylor, M. Fehr, R. Siegwart, and J. Nieto, "Voxblox: Incremental 3D Euclidean Signed Distance Fields for On-board MAV planning," in *IROS*, 2017.
- [2] L. Han, F. Gao, B. Zhou, and S. Shen, "FIESTA: Fast Incremental Euclidean Distance Fields for Online Motion Planning of Aerial Robots," in *IROS*, 2019.
- [3] Y. Pan, Y. Kompis, L. Bartolomei, R. Mascaro, C. Stachniss, and M. Chli, "Voxfield: Non-Projective Signed Distance Fields for Online Planning and 3D Reconstruction," in *IROS*, 2022.
- [4] J. J. Park, P. Florence, J. Straub, R. Newcombe, and S. Lovegrove, "DeepSDF: Learning Continuous Signed Distance Functions for Shape Representation," in *CVPR*, 2019.
- [5] T. Takikawa, J. Litalien, K. Yin, K. Kreis, C. Loop, D. Nowrouzezahrai, A. Jacobson, M. McGuire, and S. Fidler, "Neural geometric level of detail: Real-time rendering with implicit 3d shapes," in *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 11353–11362, 2021.
- [6] J. Ortiz, A. Clegg, J. Dong, E. Sucar, D. Novotny, M. Zollhoefer, and M. Mukadam, "isdf: Real-time neural signed distance fields for robot perception," in *Robotics: Science and Systems*, 2022.
- [7] B. Lee, C. Zhang, Z. Huang, and D. D. Lee, "Online continuous mapping using gaussian process implicit surfaces," in *2019 International Conference on Robotics and Automation (ICRA)*, pp. 6884–6890, 2019.
- [8] L. Wu, K. M. B. Lee, L. Liu, and T. Vidal-Calleja, "Faithful Euclidean distance field from log-Gaussian process implicit surfaces," *IEEE Robotics and Automation Letters*, vol. 6, no. 2, pp. 2461–2468, 2021.
- [9] L. Wu, C. Le Gentil, and T. Vidal-Calleja, "VDB-GPDF: Online Gaussian Process Distance Field with VDB Structure," *IEEE Robotics and Automation Letters*, 2025.
- [10] A. Robey, H. Hu, L. Lindemann, H. Zhang, D. V. Dimarogonas, S. Tu, and N. Matni, "Learning control barrier functions from expert demonstrations," in *Proc. IEEE Conf. on Decision and Control*, 2020.
- [11] K. Long, C. Qian, J. Cortés, and N. A. Atanasov, "Learning barrier functions with memory for robust safe navigation," *IEEE Robotics and Automation Letters*, vol. 6, no. 3, pp. 4931–4938, 2021.
- [12] A. Peruffo, D. Ahmed, and A. Abate, "Automated and formal synthesis of neural barrier certificates for dynamical models," *Tools and Algorithms for the Construction and Analysis of Systems*, vol. 12651, pp. 370–388, 2021.
- [13] A. Begzadić, A. Lederer, J. Cortés, and S. Herbert, "Learning high-order CBFs using Gaussian processes for systems in Brunovský canonical form," in *Proc. IEEE Conf. Decision and Control*, pp. 2945–2950, 2025.
- [14] K. Long, Y. Yi, Z. Dai, S. Herbert, J. Cortés, and N. Atanasov, "Sensor-based distributionally robust control for safe robot navigation in dynamic environments," *International Journal of Robotics Research*, 2025.
- [15] J. J. Choi, D. Lee, K. Sreenath, C. J. Tomlin, and S. L. Herbert, "Robust control barrier-value functions for safety-critical control," in *Proc. IEEE Conf. Decision and Control*, pp. 6814–6821, 2021.
- [16] S. Tonkens and S. Herbert, "Refining control barrier functions through Hamilton-Jacobi reachability," in *IEEE/RSJ Int. Conf. on Intelligent Robots & Systems*, 2022.
- [17] S. Tonkens, S. Kojima, C. Liu, J. Masri, and S. Herbert, "Refining almost-safe value functions on the fly," *arXiv preprint arXiv:2602.23478*, 2026.
- [18] A. D. Ames, X. Xu, J. W. Grizzle, and P. Tabuada, "Control barrier function based quadratic programs for safety critical systems," *IEEE Transactions on Automatic Control*, vol. 62, no. 8, pp. 3861–3876, 2017.
- [19] C. E. Rasmussen and C. K. I. Williams, *Gaussian Processes for Machine Learning*. Cambridge: MIT Press, 2006.
- [20] R. Senanayake and F. Ramos, "Bayesian hilbert maps for dynamic continuous occupancy mapping," in *Conf. on Robot Learning*, 2017.
- [21] W. E. Lorensen and H. E. Cline, "Marching cubes: A high resolution 3d surface construction algorithm," *SIGGRAPH Comput. Graph.*, vol. 21, p. 163–169, Aug. 1987.
- [22] S. R. S. Varadhan, "On the Behavior of the Fundamental Solution of the Heat Equation with Variable Coefficients," *Communications on Pure and Applied Mathematics*, 1967.
- [23] N. Srinivas, A. Krause, S. M. Kakade, and M. W. Seeger, "Information-theoretic regret bounds for Gaussian process optimization in the bandit setting," *IEEE Transactions on Information Theory*, vol. 58, p. 3250–3265, May 2012.
- [24] S. Koenig and M. Likhachev, "D* lite," in *Eighteenth national conference on Artificial intelligence*, pp. 476–483, 2002.